

* Problem 2: Gradient Descent Algo. and Logistic Regression.

(1) Likelihood function with respect to parameter w is given by,

$$L(w) = \prod_{n=1}^N p(c_i | x_n)^{y_n} (1 - p(c_i | x_n))^{1-y_n}$$

→ Applying negative logarithm on both sides,

$$\begin{aligned} \therefore -\ln(L(w)) &= -\ln \left[\prod_{n=1}^N p(c_i | x_n)^{y_n} (1 - p(c_i | x_n))^{1-y_n} \right] \\ &= - \sum_{n=1}^N \left[(y_n \ln(p(c_i | x_n)) + (1-y_n) \ln(1 - p(c_i | x_n))) \right] \end{aligned}$$

.....(i)

→ Now, $p(c_i | x) = \sigma(w^T x + w_0) = f(x)$

$\therefore p(c_i | x_n) = \sigma(w^T x_n + w_0) = f(x_n)$

.....(ii)

→ Also, $\sigma(a) = \frac{1}{1 + \exp(-a)}$ and $a(x) = w^T x + w_0$

$\therefore \sigma(a) = \frac{1}{1 + \exp(-(w^T x + w_0))}$

→ Taking derivation of above function with respect to w ,

$$\frac{\partial (\sigma(a))}{\partial w} = \frac{\partial}{\partial w} \left(\frac{1}{1 + \exp(-a)} \right)$$

$$= \frac{\partial}{\partial a} \left(\frac{1}{1 + \exp(-a)} \right) \times \frac{\partial a}{\partial w}$$

↳ (Partial derivatives)

$$= \frac{1}{(1 + \exp(-a))^2} \cdot \exp(-a) \cdot \frac{\partial a}{\partial w}$$

$$= \frac{\exp(-a)}{(1 + \exp(-a))^2}$$

$$= \frac{1}{1 + \exp(-a)} \times \frac{(1 + \exp(-a) - 1)}{1 + \exp(-a)} \cdot \frac{\partial a}{\partial w}$$

↳ Adding and Subtracting 1

$$= \frac{1}{1 + \exp(-a)} \cdot \left(1 - \frac{1}{1 + \exp(-a)} \right) \cdot \frac{\partial a}{\partial w}$$

$$= \sigma(a) (1 - \sigma(a)) \cdot \frac{\partial a}{\partial w}$$

$$\therefore \frac{\partial (\sigma(a))}{\partial w} = \sigma(a) (1 - \sigma(a)) \cdot \frac{\partial a}{\partial w}$$

(iii)

→ Now, $-\ln L(\omega) = E(\omega)$

∴ from equations (i) and (ii)

$$-\ln L(\omega) = -\sum_{n=1}^N (y_n \ln(\sigma(a)) + (1-y_n) \ln(1-\sigma(a)))$$

$$= E(\omega)$$

→ Now taking derivation of above equation with respect to ω ,

$$\frac{\partial E(\omega)}{\partial \omega} = -\sum_{n=1}^N \left(y_n \frac{\partial \ln(\sigma(a))}{\partial \omega} + (1-y_n) \frac{\partial \ln(1-\sigma(a))}{\partial \omega} \right)$$

$$= -\sum_{n=1}^N \left(y_n \frac{1}{\sigma(a)} \frac{\partial (\sigma(a))}{\partial \omega} + (1-y_n) \left(\frac{1}{\sigma(a)-1} \right) \frac{\partial (\sigma(a))}{\partial \omega} \right)$$

$$= -\sum_{n=1}^N \left(y_n \frac{1}{\sigma(a)} \frac{\partial (\sigma(a))}{\partial \omega} + (1-y_n) \left(\frac{1}{\sigma(a)-1} \right) \frac{\partial (\sigma(a))}{\partial \omega} \right)$$

→ From equation (iii) we get,

$$\frac{\partial E(\omega)}{\partial \omega} = -\sum_{n=1}^N \left(y_n \frac{1}{\sigma(a)} \sigma(a)(1-\sigma(a)) \frac{\partial a}{\partial \omega} + \frac{(1-y_n) \sigma(a)(1-\sigma(a))}{\sigma(a)-1} \frac{\partial a}{\partial \omega} \right)$$

→ According to $a(x) = w^T x + w_0$

we get $\frac{\partial a}{\partial w} = x$

∴ from above value, we get,

$$\begin{aligned} \rightarrow \frac{\partial \epsilon(w)}{\partial w} &= - \sum_{n=1}^N (y_n(1 - \sigma(a)) x_n + (y_n - 1) \cdot \sigma(a) x_n) \\ &= - \sum_{n=1}^N (y_n(1 - f(x_n)) x_n + (y_n - 1) \cdot f(x_n) x_n) \end{aligned}$$

↳ from (ii)

$$= - \sum_{n=1}^N (y_n x_n - y_n f(x_n) x_n + y_n f(x_n) x_n - f(x_n) x_n)$$

$$= - \sum_{n=1}^N x_n (y_n - f(x_n))$$

→ Thus we have obtained following result,

$$\nabla_w \epsilon(w) = \sum_{n=1}^N (f(x_n) - y_n) x_n$$

Problem 2. Gradient Descent Algorithm and Logistic Regression

Importing Libraries

In [118...

```
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import KFold
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import cross_val_score
from numpy import log, dot, e
from sklearn.model_selection import KFold
import random
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import MinMaxScaler
from sklearn.datasets import load_breast_cancer
```

Reading the dataset

In [119...

```
url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer/
headers = ['id', 'type', 'mean radius', 'mean texture', 'mean perimeter', 'me
dataset = pd.read_csv(url, names = headers)

dataset
```

Out[119...

| | id | type | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity |
|-----|----------|------|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|
| 0 | 842302 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.30010 |
| 1 | 842517 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.08690 |
| 2 | 84300903 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.19740 |
| 3 | 84348301 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.24140 |
| 4 | 84358402 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.19800 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 564 | 926424 | M | 21.56 | 22.39 | 142.00 | 1479.0 | 0.11100 | 0.11590 | 0.24390 |
| 565 | 926682 | M | 20.13 | 28.25 | 131.20 | 1261.0 | 0.09780 | 0.10340 | 0.14400 |
| 566 | 926954 | M | 16.60 | 28.08 | 108.30 | 858.1 | 0.08455 | 0.10230 | 0.09250 |
| 567 | 927241 | M | 20.60 | 29.33 | 140.10 | 1265.0 | 0.11780 | 0.27700 | 0.35140 |
| 568 | 92751 | B | 7.76 | 24.54 | 47.92 | 181.0 | 0.05263 | 0.04362 | 0.00000 |

569 rows × 32 columns

Changes to dataset for ease of use

In [120...

```

# Replacing M with 1 and B with 0
dataset = dataset.replace('M',1)
dataset = dataset.replace('B',0)

# Converting to array
y = dataset[dataset.columns[1:2]]
y = y.to_numpy().reshape(len(temp))
X = dataset[dataset.columns[2:]]
X = X.to_numpy()

dataset

```

Out[120...

| | id | type | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity |
|-----|----------|------|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|
| 0 | 842302 | 1 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.30010 |
| 1 | 842517 | 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.08690 |
| 2 | 84300903 | 1 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.19740 |
| 3 | 84348301 | 1 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.24140 |
| 4 | 84358402 | 1 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.19800 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 564 | 926424 | 1 | 21.56 | 22.39 | 142.00 | 1479.0 | 0.11100 | 0.11590 | 0.24390 |
| 565 | 926682 | 1 | 20.13 | 28.25 | 131.20 | 1261.0 | 0.09780 | 0.10340 | 0.14400 |
| 566 | 926954 | 1 | 16.60 | 28.08 | 108.30 | 858.1 | 0.08455 | 0.10230 | 0.09250 |
| 567 | 927241 | 1 | 20.60 | 29.33 | 140.10 | 1265.0 | 0.11780 | 0.27700 | 0.35140 |
| 568 | 92751 | 0 | 7.76 | 24.54 | 47.92 | 181.0 | 0.05263 | 0.04362 | 0.00000 |

569 rows × 32 columns

Stochastic Gradient Descent

In [121...

```

class LR_Stochastic:

    # Sigmoid function
    def _sigmoid(self, x):
        return 1/(1 + np.exp(-x))

    # instance variables
    def __init__(self, lr = 0.0002, iters = 10000):
        self.lr = lr
        self.iters = iters
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        # initialize parameters
        samples, features = X.shape
        # set weights to zero
        self.weights = np.zeros(features)
        self.bias = 0
        costs = []
        epochs = []

        # stochastic gradient descent algorithm

```

```

for i in range(self.iters):
    r_index = random.randint(0,samples - 1)

    sample_x = X[r_index]
    sample_y = y[r_index]
    # logistic regression equation values
    linear_model = np.dot(sample_x, self.weights) + self.bias
    # sigmoid values
    y_pred = self._sigmoid(linear_model)
    # derivatiation
    dw = (1/samples) * np.dot(sample_x.T,y_pred - sample_y)
    db = (1/samples) * np.sum(y_pred - sample_y)
    # updating the weights
    self.weights -= self.lr*dw
    self.bias -= self.lr*db
    # cost function
    cost = np.square(sample_y-y_pred)

    if i%100==0: # at every 100th iteration record the cost and iters
        costs.append(cost)
        epochs.append(i)

    return costs,epochs

# Testing the model
def predict(self, X):
    linear_model = np.dot(X, self.weights) + self.bias
    y_pred = self._sigmoid(linear_model)
    y_pred_cls = [1 if i > 0.5 else 0 for i in y_pred]
    return y_pred_cls

```

Implementing, training, testing, evaluating, recall, precision, and accuracy

In [122...

```

if __name__ == "__main__":

    # accuracy of the predictions
    def accuracy(y_true, y_pred):
        accuracy = np.sum(y_true == y_pred) / len(y_true)
        return accuracy

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,

    regressor = LR_Stochastic(lr = 0.0002, iters=11000)
    axisx,axisy = regressor.fit(X_train, y_train)
    predictions = regressor.predict(X_test)

    print("Prediction accuracy:", accuracy(y_test, predictions))

```



```

# Confusion matrix and different scores
confusionM = [[0,0],[0,0]]

for i in range(len(y_test)):
    if(y_test[i] == 0 and predictions[i] == 0):
        confusionM[0][0] += 1
    if(y_test[i] == 0 and predictions[i] == 1):
        confusionM[0][1] += 1
    if(y_test[i] == 1 and predictions[i] == 0):
        confusionM[1][0] += 1
    if(y_test[i] == 1 and predictions[i] == 1):
        confusionM[1][1] += 1

# true negative, false positive, false negative, true positive
tn = confusionM[0][0]
fp = confusionM[0][1]
fn = confusionM[1][0]
tp = confusionM[1][1]

precision_score = tp/(fp + tp)
recall_score = tp/(fn + tp)
accuracy_score = (tp + tn)/(tp + fn + tn + fp)
print('Precision: %.3f' % precision_score)
print('Recall: %.3f' % recall_score)
print('Accuracy: %.3f' % accuracy_score)

plt.xlabel("epoch")
plt.ylabel("cost")
plt.plot(axisy,axisx)

#cross-validation
kFold = KFold(n_splits=3, random_state=None)

acc_score = []

for train , test in kFold.split(X):
    X_train , X_test = X[train,:],X[test,:]
    y_train , y_test = y[train] , y[test]

    regressor.fit(X_train,y_train)
    pred_values = regressor.predict(X_test)


    acc = accuracy(pred_values, y_test)
    acc_score.append(acc)

avg_acc_score = sum(acc_score)/3

print('accuracy of each fold - {}'.format(acc_score))
print('Avg accuracy : {}'.format(avg_acc_score))

```

Prediction accuracy: 0.8951048951048951
 Precision: 0.975
 Recall: 0.736
 Accuracy: 0.895
 accuracy of each fold - [0.9157894736842105, 0.8894736842105263, 0.9312169312169312]
 Avg accuracy : 0.9121600297038893



Mini-Batch gradient descent

In [123...

```

class LR_Mini_Batch:

    # Sigmoid function
    def _sigmoid(self, x):
        return 1/(1 + np.exp(-x))

    # defining instance variables
    def __init__(self, lr = 0.0001, iters = 200, batch_size = 20):
        self.lr = lr
        self.iters = iters
        self.weights = None
        self.bias = None
        self.batch_size = batch_size

    def fit(self, X, y):
        # initialize parameters
        samples, features = X.shape
        # set all weights equal to 0 and bias equal to 0
        self.weights = np.zeros(features)
        self.bias = 0
        costs = []
        epochs = []

        # mini-batch gradient descent algorithm
        for i in range(self.iters):
  
```

```

r_indices = np.random.permutation(samples)
sample_x = X[r_indices]
sample_y = y[r_indices]

# iterations in batches
for j in range(0,samples,self.batch_size):
    Xt = sample_x[j:j+self.batch_size]
    yt = sample_y[j:j+self.batch_size]

    # linear logistic regression equation
    linear_model = np.dot(Xt, self.weights) + self.bias

    # sigmoid values
    y_pred = self._sigmoid(linear_model)

    # derivativation
    dw = (1/samples) * np.dot(Xt.T,y_pred - yt)
    db = (1/samples) * np.sum(y_pred - yt)

    # updating the weights
    self.weights -= self.lr*dw
    self.bias -= self.lr*db

    # MeanSquaredError of costs
    cost = np.mean(np.square(yt-y_pred))

    if i%10==0: # at every 10th iteration record the cost and iters v
        costs.append(cost)
        epochs.append(i)

    return costs,epochs

# Testing of model
def predict(self, X):
    linear_model = np.dot(X, self.weights) + self.bias
    y_pred = self._sigmoid(linear_model)
    y_pred_cls = [1 if i > 0.5 else 0 for i in y_pred]
    return y_pred_cls

```

Implementing, training, testing, evaluating, recall, precision, and accuracy

In [124...

```

if __name__ == "__main__":

    # accuracy of predictions
    def accuracy(y_true, y_pred):
        accuracy = np.sum(y_true == y_pred) / len(y_true)
        return accuracy

```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,

regressor = LR_Mini_Batch(lr = 0.0001, iters=100, batch_size = 10)
axisx,axisy = regressor.fit(X_train, y_train)
predictions = regressor.predict(X_test)

print("Prediction accuracy:", accuracy(y_test, predictions))

plt.xlabel("epoch")
plt.ylabel("cost")
plt.plot(axisy,axisx)

# Confusion matrix and different scores
confusionM = [[0,0],[0,0]]

for i in range(len(y_test)):
    if(y_test[i] == 0 and predictions[i] == 0):
        confusionM[0][0] += 1
    if(y_test[i] == 0 and predictions[i] == 1):
        confusionM[0][1] += 1
    if(y_test[i] == 1 and predictions[i] == 0):
        confusionM[1][0] += 1
    if(y_test[i] == 1 and predictions[i] == 1):
        confusionM[1][1] += 1

# true negative, false positive, false negative, true positive
tn = confusionM[0][0]
fp = confusionM[0][1]
fn = confusionM[1][0]
tp = confusionM[1][1]

precision_score = tp/(fp + tp)
recall_score = tp/(fn + tp)
accuracy_score = (tp + tn)/(tp + fn + tn + fp)
print('Precision: %.3f' % precision_score)
print('Recall: %.3f' % recall_score)
print('Accuracy: %.3f' % accuracy_score)

# cross-validation
kFold = KFold(n_splits=3, random_state=None)

acc_score = []

for train , test in kFold.split(X):
    X_train , X_test = X[train,:],X[test,:]
    y_train , y_test = y[train] , y[test]

    regressor.fit(X_train,y_train)
    pred_values = regressor.predict(X_test)

    acc = accuracy(pred_values , y_test)
    acc_score.append(acc)
```



```
avg_acc_score = sum(acc_score)/3

print('accuracy of each fold - {}'.format(acc_score))
print('Avg accuracy : {}'.format(avg_acc_score))
```

Prediction accuracy: 0.8951048951048951

Precision: 0.975

Recall: 0.736

Accuracy: 0.895

accuracy of each fold - [0.8894736842105263, 0.9052631578947369, 0.9259259259259259]

Avg accuracy : 0.9068875893437297

